

HW1-jy2913

February 12, 2019

Sole team member team: Name: Jin Yan UNI: jy2913 Link to Colaboratory:
https://colab.research.google.com/drive/1LFp0QTjDMMv0wDrZK9CqP4pFcUrZGIzs#scrollTo=qsH_VY6Aqw

```
In [0]: #download CIFAR10 dataset from pytorch
import torch, torchvision
transform = torchvision.transforms.Compose([torchvision.transforms.ToTensor(),torchvisi
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True,trans
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, trans
```

Files already downloaded and verified
Files already downloaded and verified

```
In [0]: import numpy as np
import matplotlib.pyplot as plt
import scipy.misc
```

```
In [0]: """ Pre-processing the dataset """
# get X input from trainset
X_train = trainset.train_data
print(X_train.shape)
# flatten (50000,32,32,3) size of X into (3072,50000) size of matrix
# normalize each pixel value to between 0 and 1
X_train = X_train.reshape(X_train.shape[0],-1).T / 255
print(X_train.shape)
# did the same thing on testset: got X input and flatten it
X_test = testset.test_data
print(X_test.shape)
X_test = X_test.reshape(X_test.shape[0],-1).T / 255
print(X_test.shape)
# get label y from trainset
y_train = trainset.train_labels
print(y_train)
# reshape y from list to numpy array
y_train = np.reshape(np.asarray(y_train),(50000,))
print(y_train.shape)
# did the same thing on testset: got y label and reshape it
y_test = testset.test_labels
```

```

y_test = np.reshape(np.asarray(y_test),(10000,))
print(y_test.shape)

# function takes X input and y labels and ratio to split validation set
def split(X, y, val_ratio):
    val_number = int(val_ratio * X.shape[1])
    random_indice = np.random.permutation(X.shape[1])
    return X[:, random_indice[val_number:]], y[random_indice[val_number:]], X[:, random_indice[:val_number]]
# split validation set from training set
X_train, y_train, X_val, y_val = split(X_train, y_train, val_ratio = 0.1)

(50000, 32, 32, 3)
(3072, 50000)
(10000, 32, 32, 3)
(3072, 10000)
[6, 9, 9, 4, 1, 1, 2, 7, 8, 3, 4, 7, 7, 2, 9, 9, 9, 3, 2, 6, 4, 3, 6, 6, 2, 6, 3, 5, 4, 0, 0, 9]
(50000,)
(10000,)

In [0]: class NeuralNetwork():

    def __init__(self, layer_dims):
        """
        Arguments:
            layer_dims -- A list contains the dimensions of each layer in CNN.

        Attributes generated:
            parameters -- a dict contains parameters "W1", "b1", ..., "WL", "bL" of each convolutional layer
                Wl -- weight matrix of shape (layer_dims[l], layer_dims[l-1])
                bl -- bias vector of shape (layer_dims[l], 1)
            num_layers -- the length of layer_dims list
        """

        np.random.seed(1)
        self.num_layers = len(layer_dims)
        self.layer_dims = layer_dims
        self.parameters = {}
        L = len(self.layer_dims)
        for l in range(1, L):
            self.parameters['W'+ str(l)] = np.random.randn(layer_dims[l], layer_dims[l-1])
            self.parameters['b'+ str(l)] = np.zeros([layer_dims[l], 1])
        self.X_val = None
        self.y_val = None

    def affineForward(self, A, W, b):
        """

```

Implement the linear portion of CNN's forward propagation.

Arguments:

A -- activation from previous layer

W -- weights matrix

b -- bias vector

Returns:

Z -- the input of the activation function

cache -- a dict stores "A", "W" and "b" during forward propagation

"""

Z = np.matmul(W, A) + b

cache = (A, W, b, Z)

return Z, cache

```
def activationForward(self, Z, activation="relu"):
```

"""

Implement the linear to activation portion of CNN's forward propagation

Arguments:

Z -- output from linear portion of forward propagation

activation(type) -- "relu" in this case

Returns:

A -- the output of relu activation

cache -- a dict stores "A", "W" and "b" during forward propagation

"""

A = np.maximum(0, Z)

assert (Z.shape == A.shape)

return A

```
def forwardPropagation(self, X):
```

"""

Implement the forward propagation

Arguments:

X -- input from input layer as the starting point of forward propagation

Returns:

A -- the output of the activation function for each layer

caches -- a list stores parameters (cache) for each layer during forward propagation

"""

A = X

caches = []

for l in range(1, self.num_layers):

 Z, cache = self.affineForward(A, self.parameters['W'+ str(l)], self.parameters['b'+ str(l)])

 caches.append(cache)

```

        if (l <= self.num_layers - 1):
            A = self.activationForward(Z)
    return A, caches

def softmax(self,AL):
    """
    Implement the softmax function

    Arguments:
    AL -- output of linear portion of the output layer (last layer)

    Returns:
    p -- softmax probability
    """
    m = AL.shape[1]
    p = np.exp(AL - np.max(AL, axis=0, keepdims=True))
    p /= np.sum(p, axis=0, keepdims=True)
    return p

def costFunction(self, AL, y):
    """
    Implement the cross entropy loss

    Arguments:
    AL -- output of linear portion of the output layer (last layer)
    y -- the labels of data

    Returns:
    cost -- the cross entropy loss value of each iteration of forward propagation
    """
    m = AL.shape[1]
    p = self.softmax(AL)
    cost = -np.sum(np.log(p[y, np.arange(m)])) / m
    return cost

def derivative_cost(self, AL, y):
    """
    Implement the first step of back propagation: the derivative of cost function

    Arguments:
    AL -- output of linear portion of the output layer (last layer)
    y -- the labels of data

    Returns:
    dAL -- the derivative of cost function over AL

```

```

        """
        m = AL.shape[1]
        p = self.softmax(AL)
        dAL = p.copy()
        dAL[y, np.arange(m)] -= 1
        dAL /= m
        return dAL

def affineBackward(self, dA_prev, cache):
    """
    Implement the linear portion of backward propagation of one layer (layer l)

    Arguments:
        dAL -- derivative of the cost with respect to the activation output (of current layer)
        cache -- releases A, W, b, Z values stored from the tuple during forward propagation

    Returns:
        dA -- Gradient of the cost over activation output from previous layer (layer l-1)
        dW -- Gradient of the cost over W for the current layer(layerl)
        db -- Gradient of the cost over b for the current layer(layerl)
    """
    m = A.shape[1]
    A, W, b, Z = cache
    dA = np.matmul(W.T, dA_prev)
    dW = np.matmul(dA_prev, A.T) / m
    db = np.sum(dA_prev, axis = 1, keepdims=True) / m
    return dA, dW, db

def activationBackward(self, dA, cache, activation="relu"):
    """
    Implement the derivative of cost function over relu activation input Z

    Arguments:
        dA -- the gradient of cost function over relu activation output A
        cache -- a tuple stores A,W,b,Z for each layer, cache[3] is Z
        activation(type) -- 'relu'

    Returns:
        relu_backward -- the derivative of cost function over Z
    """

    relu_backward = self.derivative_relu(dA, cache[3])
    return relu_backward

def derivative_relu(self, dA, cache):
    """
    Implement the derivative calculation of relu activation during backpropagation

```

```

Arguments:
dA -- the gradient of cost function over relu activation output A
cache -- a tuple stores A,W,b,Z for each layer, cache[3] is Z

Returns:
dZ -- the derivative of cost function over Z
"""

calcu = np.maximum(0, cache)
calcu[out > 0] = 1
dZ = calcu * dA
return dZ

def backPropagation(self, dAL, y, caches):
    """
    Implement backpropagation for each layer

    Arguments:
    dAL -- the gradient of cost function over the activation output of last layer:
    y -- the labels of data
    caches -- a list stores tuple of A,W,b,Z for each layer

    Returns:
    grads -- a dict stores dWl and dbl for each layer l
    """
    dA = dAL
    grads = {}
    for l in reversed(range(0, self.num_layers-1)):
        if l < self.num_layers - 1:
            dA = self.activationBackward(dA, caches[l-1])
            dA, dW, db = self.affineBackward(dA, caches[l-1])
            grads['W'+ str(l)] = dW
            grads['b'+ str(l)] = db
    return grads

def updateParameters(self, grads, alpha):
    """
    Use gradient descent to implement parameters update

    Arguments:
    grads -- a dict stores all parameters gradients for each layer
    alpha -- learning rate

    Returns:
    self.parameters -- a dict stores updated parameters for W and b of each layer
        parameters["W" + str(l)] = ...
        parameters["b" + str(l)] = ...
    """
    L = len(self.parameters) // 2

```

```

for l in range(1,L-1):
    self.parameters["W" + str(l)] -= alpha*grads["W"+str(l)]
    self.parameters["b" + str(l)] -= alpha*grads["b"+str(l)]
return self.parameters

def normalized_X(self,X):
    """
    Standardize the input data by subtract mean and divided by variance

    Arguments:
    X -- input data

    Return:
    norm_X -- standardized data
    """

    standardize = {}
    standardize['mean'] = np.mean(X, axis = 1, keepdims = True)
    standardize['var'] = np.var(X, axis = 1, keepdims = True)
    norm_X = (X - standardize['mean']) / np.sqrt(standardize['var'])
    return norm_X

def train(self, X, y, iters, alpha, batch_size, print_every):
    """
    It takes advantage of every function in this class to implement training and validation

    Arguments:
    X -- input data
    y -- labels of data
    iters -- number of iterations to run
    alpha -- learning rate
    batch_size -- number of samples to assign to minibatch
    print_every -- number of iterations to print

    Return:
    Number of iterations, train loss, train_acc, and valid_acc in every 100 iterations

    """

    X = self.normalized_X(X)
    for i in range(0, iters):
        X_batch, y_batch = self.get_batch(X, y, batch_size)
        AL, cache = self.forwardPropagation(X_batch)
        loss = self.costFunction(AL, y_batch)
        dAL = self.derivative_cost(AL, y_batch)
        grads = self.backPropagation(dAL, y_batch, cache)
        self.updateParameters(grads, alpha)

```

```

        if i % print_every == 0:
            train_acc = self.score(self.predict(X), y)
            val_acc = self.score(self.predict(self.X_val), self.y_val)
            print('iter={:5}, loss={:.4f}, train_acc={:.4f}, validation_acc={:.4f}')

def predict(self, X):
    """
    It predicts the label given . input x
    Argument:
    X -- input of data
    Return:
    y_pred -- predicted label
    """
    X = self.normalized_X(X)
    AL, _ = self.forwardPropagation(X)
    y_pred = np.argmax(AL, axis = 0)
    return y_pred

def score(self, y_pred, y_true):
    """
    It calculates the percentage of correct predicted labels over true labels

    Argument:
    y_pred-- predicted labels
    y_true -- true labels

    Return: percentage of correct predicted labels
    """
    correct = np.mean(y_pred == y_true)
    return correct

def load_validation_set(self, X_val, y_val):
    """
    Load validation set to CNN
    """
    self.X_val = X_val
    self.y_val = y_val

def get_batch(self, X, y, batch_size):
    """
    Load minibatch to CNN
    """

    batch_index = np.random.randint(X.shape[1], size = batch_size)
    X_batch = X[:, batch_index]
    y_batch = y[batch_index]

```



```
return X_batch, y_batch
```

```
In [0]: """
```

```
Train CNN and validate at the same time
```

```
"""
```

```
# I tried idifferent architecture of CNN and it turns out 3 hidden layers of 1024, 256, 128
```

```
# I also tried different learning drate and batch size and it turns out alpha=1 and batch_size=100
```

```
# In this trained CNN, I reach 54.12% validation accuracy and 52.99% test accuracy in
```

```
layer_dims = [X_train.shape[0], 1024, 256, 128, 10]
```

```
CNN = NeuralNetwork(layer_dims)
```

```
CNN.load_validation_set(X_val, y_val)
```

```
CNN.train(X_train, y_train, iters=10000, alpha=1, batch_size=100, print_every=100)
```

```
iter=    0, loss=2.3617, train_acc=0.0999, validation_acc=0.1064
iter=   100, loss=2.0751, train_acc=0.2959, validation_acc=0.2948
iter=   200, loss=1.8461, train_acc=0.3493, validation_acc=0.3530
iter=   300, loss=1.7975, train_acc=0.3786, validation_acc=0.3762
iter=   400, loss=1.6927, train_acc=0.4006, validation_acc=0.4002
iter=   500, loss=1.7257, train_acc=0.4153, validation_acc=0.4088
iter=   600, loss=1.6456, train_acc=0.4278, validation_acc=0.4176
iter=   700, loss=1.7837, train_acc=0.4386, validation_acc=0.4272
iter=   800, loss=1.7971, train_acc=0.4468, validation_acc=0.4314
iter=   900, loss=1.6531, train_acc=0.4553, validation_acc=0.4396
iter=  1000, loss=1.6428, train_acc=0.4612, validation_acc=0.4514
iter=  1100, loss=1.5892, train_acc=0.4674, validation_acc=0.4504
iter=  1200, loss=1.6917, train_acc=0.4736, validation_acc=0.4532
iter=  1300, loss=1.4878, train_acc=0.4779, validation_acc=0.4560
iter=  1400, loss=1.4416, train_acc=0.4837, validation_acc=0.4680
iter=  1500, loss=1.4070, train_acc=0.4885, validation_acc=0.4596
iter=  1600, loss=1.3062, train_acc=0.4958, validation_acc=0.4642
iter=  1700, loss=1.4820, train_acc=0.4989, validation_acc=0.4618
iter=  1800, loss=1.3706, train_acc=0.5010, validation_acc=0.4672
iter=  1900, loss=1.4472, train_acc=0.5081, validation_acc=0.4714
iter=  2000, loss=1.6072, train_acc=0.5118, validation_acc=0.4750
iter=  2100, loss=1.4100, train_acc=0.5090, validation_acc=0.4700
iter=  2200, loss=1.4180, train_acc=0.5184, validation_acc=0.4834
iter=  2300, loss=1.4381, train_acc=0.5197, validation_acc=0.4812
iter=  2400, loss=1.4284, train_acc=0.5244, validation_acc=0.4872
iter=  2500, loss=1.3580, train_acc=0.5291, validation_acc=0.4840
iter=  2600, loss=1.2492, train_acc=0.5353, validation_acc=0.4894
iter=  2700, loss=1.5415, train_acc=0.5384, validation_acc=0.4938
iter=  2800, loss=1.4538, train_acc=0.5415, validation_acc=0.4924
iter=  2900, loss=1.2612, train_acc=0.5461, validation_acc=0.4922
iter=  3000, loss=1.3019, train_acc=0.5513, validation_acc=0.4980
iter=  3100, loss=1.2984, train_acc=0.5501, validation_acc=0.5016
iter=  3200, loss=1.1218, train_acc=0.5572, validation_acc=0.4976
iter=  3300, loss=1.3844, train_acc=0.5569, validation_acc=0.5042
```

iter= 3400, loss=1.3159, train_acc=0.5593, validation_acc=0.5064
iter= 3500, loss=1.3674, train_acc=0.5641, validation_acc=0.5018
iter= 3600, loss=1.3703, train_acc=0.5645, validation_acc=0.5016
iter= 3700, loss=1.4296, train_acc=0.5690, validation_acc=0.5096
iter= 3800, loss=1.4543, train_acc=0.5718, validation_acc=0.5084
iter= 3900, loss=1.2558, train_acc=0.5754, validation_acc=0.5092
iter= 4000, loss=1.2729, train_acc=0.5790, validation_acc=0.5084
iter= 4100, loss=1.1926, train_acc=0.5764, validation_acc=0.5186
iter= 4200, loss=1.3360, train_acc=0.5815, validation_acc=0.5102
iter= 4300, loss=1.1685, train_acc=0.5840, validation_acc=0.5126
iter= 4400, loss=1.0960, train_acc=0.5850, validation_acc=0.5154
iter= 4500, loss=1.2709, train_acc=0.5880, validation_acc=0.5202
iter= 4600, loss=1.1889, train_acc=0.5892, validation_acc=0.5116
iter= 4700, loss=1.2040, train_acc=0.5927, validation_acc=0.5222
iter= 4800, loss=1.1178, train_acc=0.5971, validation_acc=0.5148
iter= 4900, loss=1.0266, train_acc=0.6021, validation_acc=0.5226
iter= 5000, loss=1.3410, train_acc=0.6036, validation_acc=0.5230
iter= 5100, loss=1.1866, train_acc=0.6079, validation_acc=0.5294
iter= 5200, loss=1.0812, train_acc=0.6103, validation_acc=0.5268
iter= 5300, loss=1.1771, train_acc=0.6108, validation_acc=0.5256
iter= 5400, loss=1.3337, train_acc=0.6172, validation_acc=0.5304
iter= 5500, loss=1.1678, train_acc=0.6180, validation_acc=0.5302
iter= 5600, loss=1.2106, train_acc=0.6181, validation_acc=0.5242
iter= 5700, loss=0.8486, train_acc=0.6255, validation_acc=0.5272
iter= 5800, loss=1.0979, train_acc=0.6220, validation_acc=0.5260
iter= 5900, loss=0.9381, train_acc=0.6273, validation_acc=0.5298
iter= 6000, loss=1.1777, train_acc=0.6294, validation_acc=0.5304
iter= 6100, loss=1.1077, train_acc=0.6285, validation_acc=0.5290
iter= 6200, loss=1.2029, train_acc=0.6188, validation_acc=0.5264
iter= 6300, loss=1.1024, train_acc=0.6365, validation_acc=0.5316
iter= 6400, loss=0.9901, train_acc=0.6314, validation_acc=0.5316
iter= 6500, loss=1.0325, train_acc=0.6378, validation_acc=0.5340
iter= 6600, loss=1.0798, train_acc=0.6381, validation_acc=0.5304
iter= 6700, loss=0.9387, train_acc=0.6422, validation_acc=0.5296
iter= 6800, loss=1.2228, train_acc=0.6452, validation_acc=0.5368
iter= 6900, loss=1.1876, train_acc=0.6407, validation_acc=0.5322
iter= 7000, loss=0.9095, train_acc=0.6509, validation_acc=0.5354
iter= 7100, loss=1.0665, train_acc=0.6489, validation_acc=0.5364
iter= 7200, loss=1.2499, train_acc=0.6576, validation_acc=0.5372
iter= 7300, loss=1.0148, train_acc=0.6551, validation_acc=0.5394
iter= 7400, loss=1.0043, train_acc=0.6538, validation_acc=0.5356
iter= 7500, loss=1.0910, train_acc=0.6589, validation_acc=0.5412
iter= 7600, loss=0.9433, train_acc=0.6632, validation_acc=0.5458
iter= 7700, loss=1.0436, train_acc=0.6630, validation_acc=0.5364
iter= 7800, loss=0.8558, train_acc=0.6640, validation_acc=0.5358
iter= 7900, loss=0.9926, train_acc=0.6714, validation_acc=0.5422
iter= 8000, loss=1.1019, train_acc=0.6728, validation_acc=0.5422
iter= 8100, loss=0.9499, train_acc=0.6720, validation_acc=0.5338

```

iter= 8200, loss=0.9922, train_acc=0.6770, validation_acc=0.5386
iter= 8300, loss=0.8384, train_acc=0.6788, validation_acc=0.5456
iter= 8400, loss=1.0024, train_acc=0.6706, validation_acc=0.5364
iter= 8500, loss=0.8492, train_acc=0.6794, validation_acc=0.5352
iter= 8600, loss=0.9987, train_acc=0.6816, validation_acc=0.5416
iter= 8700, loss=0.9400, train_acc=0.6886, validation_acc=0.5458
iter= 8800, loss=0.9616, train_acc=0.6854, validation_acc=0.5412
iter= 8900, loss=1.0930, train_acc=0.6908, validation_acc=0.5464
iter= 9000, loss=0.9944, train_acc=0.6914, validation_acc=0.5452
iter= 9100, loss=0.8386, train_acc=0.6916, validation_acc=0.5486
iter= 9200, loss=0.8467, train_acc=0.6893, validation_acc=0.5408
iter= 9300, loss=0.9941, train_acc=0.6972, validation_acc=0.5392
iter= 9400, loss=1.0225, train_acc=0.6920, validation_acc=0.5326
iter= 9500, loss=0.8249, train_acc=0.7022, validation_acc=0.5426
iter= 9600, loss=1.0521, train_acc=0.7055, validation_acc=0.5456
iter= 9700, loss=0.8433, train_acc=0.7016, validation_acc=0.5426
iter= 9800, loss=0.9047, train_acc=0.7047, validation_acc=0.5394
iter= 9900, loss=0.8353, train_acc=0.7116, validation_acc=0.5412

```

```

In [0]: #test set accuracy
        y_pred = CNN.predict(X_test)
        test_acc = CNN.score(y_pred, y_test)
        print('test_acc ={:4}'.format(test_acc))

```

```
test_acc =0.5299
```

```

In [0]: #visualize the prediction result of trained CNN on some of test set images
import torchvision
import torchvision.transforms as transforms
testloader = torch.utils.data.DataLoader(testset, batch_size=10,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()


```

```

# get some random training images
dataiter = iter(testloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print("testlabels: " + ' '.join('%5s' % classes[labels[j]] for j in range(10)))
print("predlabels: " + ' '.join('%5s' % classes[y_pred[j]] for j in range(10)))

```



```

testlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
predlabels:  cat  ship plane plane  deer  frog  cat  frog  dog  car

```

```

In [0]: import torch
import torchvision
import torchvision.transforms as transforms
from torchvision import datasets
import numpy as np
from torch.utils.data.sampler import SubsetRandomSampler

In [0]: #download CIFAR10 data from pytorch

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

train_dataset = datasets.CIFAR10(root='./data', train=True,
                                  download=True, transform=transform)
valid_dataset = datasets.CIFAR10(root='./data', train=True,
                                   download=True, transform=transform)

```

```

#generate index for validation set
#use SubsetRandomSampler to split validation set
valid_size = 0.1
num_train = len(train_dataset)
indices = list(range(num_train))
split = int(np.floor(valid_size * num_train))
np.random.seed(0)
np.random.shuffle(indices)

train_idx, valid_idx = indices[split:], indices[:split]
train_sampler = SubsetRandomSampler(train_idx)
valid_sampler = SubsetRandomSampler(valid_idx)

#load train set, validation set, and test set

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=50,
                                           sampler=train_sampler,num_workers=2, pin_memory=True)
valid_loader = torch.utils.data.DataLoader(train_dataset, batch_size=50,
                                           sampler=valid_sampler,num_workers=2, pin_memory=True)
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=50,
                                         shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz
 Files already downloaded and verified
 Files already downloaded and verified

```

In [0]: import matplotlib.pyplot as plt
import numpy as np

```

```

def imshow(img):
    img = img / 2 + 0.5
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

```

```

dataiter = iter(train_loader)
images, labels = dataiter.next()

```

```
imshow(torchvision.utils.make_grid(images))

print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



horse dog bird frog

```
In [0]: import torch.nn as nn
import torch.nn.functional as F
# model_b is my baseline model. All my following models are generated upon this model
# model_b consists of two convolutional layers and two fc layers
# model_b has validation accuracy and test accuracy of 73%

class CIFAR10_base(nn.Module):
    def __init__(self):
        super(CIFAR10_base, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 5, padding=2)
        self.conv2 = nn.Conv2d(32, 64, 5, padding=2)
        self.fc1 = nn.Linear(64*8*8, 1024)
        self.fc2 = nn.Linear(1024,10)
```

```

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, 64*8*8) # reshape before sending to fc layer
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return F.log_softmax(x)

model_b = CIFAR10_base()

In [0]: print(model_b)

CIFAR10_base(
  (conv1): Conv2d(3, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (conv2): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (fc1): Linear(in_features=4096, out_features=1024, bias=True)
  (fc2): Linear(in_features=1024, out_features=10, bias=True)
)

In [0]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model_b.parameters(), lr=0.01, momentum=0.9)

In [0]: def train(model, train_loader, criterion, optimizer, epoch):
    model.cuda()
    i = 0

    model.train() # to set train mode for drop out
    train_loss, train_accu = [], []
    for images, labels in train_loader:
        # send tensors to GPU
        images, labels = images.cuda(), labels.cuda()

        optimizer.zero_grad() # zero the parameter gradients
        outputs = model(images) # calls the forward function of model, i.e.
        loss = F.nll_loss(outputs, labels) # calculate loss
        loss.backward() # calculate gradients
        train_loss.append(loss.item())
        optimizer.step() # update learnable parameters
        predictions = outputs.data.max(1)[1] # column at idx 1 has actual prob.

    accuracy = np.sum(predictions.cpu().numpy()==labels.cpu().numpy())/batch_size*
    #return loss.item(), accuracy
    train_accu.append(accuracy)
    if i % 1000 == 0:
        print('Train Step: {} \t Train Loss: {:.3f} \t Train Accuracy: {:.3f}'.format(
            i += 1

```

```
In [0]: def validation(model, valid_loader, epoch):
        i = 0
        model.eval()
        correct = 0
        for images, labels in valid_loader:
            with torch.no_grad(): # so that computation graph history is not stored
                images, labels = images.cuda(), labels.cuda() # send tensors to GPU
                outputs = model(images)
                predictions = outputs.data.max(1)[1]
                correct += predictions.eq(labels.data).sum()
            accuracy = 100.0 * correct / (len(valid_loader.dataset) * valid_size)
            if i == 99:
                print('Validation step: {} \t Validation accuracy: {:.3f}'.format(epoch,
                    accuracy))
            i += 1
```

```
In [0]: batch_size = 50
        # send model to GPU

        #train_loss, train_accu, valid_accu = [], [], []

        for epoch in range(20):
            train(model_b, train_loader, criterion, optimizer, epoch)
            validation(model_b, valid_loader, epoch)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:19: UserWarning: Implicit dimension

Train Step: 0	Train Loss: 1.475	Train Accuracy: 46.000
Validation step: 0	Validation accuracy: 67.000	
Train Step: 1000	Train Loss: 0.680	Train Accuracy: 78.000
Validation step: 1000	Validation accuracy: 70.000	
Train Step: 2000	Train Loss: 0.768	Train Accuracy: 78.000
Validation step: 2000	Validation accuracy: 72.000	
Train Step: 3000	Train Loss: 0.296	Train Accuracy: 90.000
Validation step: 3000	Validation accuracy: 72.000	
Train Step: 4000	Train Loss: 0.227	Train Accuracy: 94.000
Validation step: 4000	Validation accuracy: 73.000	
Train Step: 5000	Train Loss: 0.195	Train Accuracy: 92.000
Validation step: 5000	Validation accuracy: 72.000	
Train Step: 6000	Train Loss: 0.091	Train Accuracy: 98.000
Validation step: 6000	Validation accuracy: 73.000	
Train Step: 7000	Train Loss: 0.081	Train Accuracy: 94.000
Validation step: 7000	Validation accuracy: 73.000	

Train Step: 8000

Train Loss: 0.012

Train Accuracy: 100.000

KeyboardInterrupt

Traceback (most recent call last)

```
<ipython-input-16-bec414c5a960> in <module>()
    5
    6 for epoch in range(20):
----> 7     train(model_b,train_loader,criterion,optimizer,epoch)
    8     validation(model_b,valid_loader,epoch)
    9

<ipython-input-10-fdc09cd7c551> in train(model, train_loader, criterion, optimizer, ep
   17     predictions = outputs.data.max(1)[1]# column at idx 1 has actual prob.
   18
---> 19     accuracy = np.sum(predictions.cpu().numpy()==labels.cpu().numpy())/batch_s
   20     #return loss.item(), accuracy
   21     train_accu.append(accuracy)
```

KeyboardInterrupt:

```
In [0]: model_b.eval()
        correct = 0
        for images, labels in testloader:
            with torch.no_grad(): # so that computation graph history is not stored
                images, labels = images.cuda(), labels.cuda() # send tensors to GPU
                outputs = model_b(images)
                predictions_T = outputs.data.max(1)[1]
                correct += predictions_T.eq(labels.data).sum()

        print('Test set accuracy: {:.2f}%'.format(100.0 * correct / len(testloader.dataset)))
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:19: UserWarning: Implicit dimension

Test set accuracy: 73.00%

```
In [0]: import torchvision
        import torchvision.transforms as transforms
        testloader = torch.utils.data.DataLoader(testset, batch_size=10,
                                                    shuffle=False, num_workers=2)
```

```
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

```
import matplotlib.pyplot as plt
import numpy as np
```


```
# functions to show an image
```

```
def imshow(img):
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

```
# get some random training images
dataiter = iter(testloader)
images, labels = dataiter.next()
```

```
# show images
imshow(torchvision.utils.make_grid(images))
images, labels = images.cuda(), labels.cuda() # send tensors to GPU
outputs = model_b(images)
predictions_T = outputs.data.max(1)[1]
correct += predictions_T.eq(labels.data).sum()
```

```
# print labels
print("testlabels: " + ' '.join('%5s' % classes[labels.data[j]] for j in range(10)))
print("predlabels: " + ' '.join('%5s' % classes[predictions_T[j]] for j in range(10)))
```



```
testlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
predlabels:  cat  car  ship plane  deer  frog  car  frog  cat  car
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:19: UserWarning: Implicit dimensi

```
In [0]: import torch.nn as nn
import torch.nn.functional as F
# model1 add a dropout layer between fc1 and fc2 of model_b
# improve val acc to 74% and test acc to 75% compared to model_b

class CIFAR10_1(nn.Module):
    def __init__(self):
        super(CIFAR10_1, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 5, padding=2)
        self.conv2 = nn.Conv2d(32, 64, 5, padding=2)
        self.fc1 = nn.Linear(64*8*8, 1024)
        self.fc2 = nn.Linear(1024,10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, 64*8*8) # reshape before sending to fc layer
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x)
        return F.log_softmax(x)

model1 = CIFAR10_1()
```

```
In [0]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model1.parameters(), lr=0.01, momentum=0.9)
```

```
In [0]: batch_size = 50
# send model to GPU

#train_loss, train_accu, valid_accu = [], [], []

for epoch in range(20):
    train(model1,train_loader,criterion,optimizer,epoch)
    validation(model1,valid_loader,epoch)
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:21: UserWarning: Implicit dimensi

Train Step: 0	Train Loss: 2.314	Train Accuracy: 6.000
Validation step: 0	Validation accuracy: 57.000	
Train Step: 1000	Train Loss: 1.200	Train Accuracy: 58.000
Validation step: 1000	Validation accuracy: 65.000	

Train Step: 2000	Train Loss: 1.041	Train Accuracy: 66.000
Validation step: 2000	Validation accuracy: 69.000	
Train Step: 3000	Train Loss: 0.485	Train Accuracy: 84.000
Validation step: 3000	Validation accuracy: 71.000	
Train Step: 4000	Train Loss: 0.633	Train Accuracy: 76.000
Validation step: 4000	Validation accuracy: 72.000	
Train Step: 5000	Train Loss: 0.478	Train Accuracy: 92.000
Validation step: 5000	Validation accuracy: 74.000	
Train Step: 6000	Train Loss: 0.571	Train Accuracy: 88.000
Validation step: 6000	Validation accuracy: 75.000	
Train Step: 7000	Train Loss: 0.204	Train Accuracy: 90.000
Validation step: 7000	Validation accuracy: 73.000	
Train Step: 8000	Train Loss: 0.329	Train Accuracy: 90.000
Validation step: 8000	Validation accuracy: 75.000	
Train Step: 9000	Train Loss: 0.129	Train Accuracy: 98.000
Validation step: 9000	Validation accuracy: 74.000	
Train Step: 10000	Train Loss: 0.130	Train Accuracy: 96.000
Validation step: 10000	Validation accuracy: 75.000	
Train Step: 11000	Train Loss: 0.103	Train Accuracy: 96.000
Validation step: 11000	Validation accuracy: 75.000	
Train Step: 12000	Train Loss: 0.075	Train Accuracy: 98.000
Validation step: 12000	Validation accuracy: 75.000	
Train Step: 13000	Train Loss: 0.148	Train Accuracy: 94.000
Validation step: 13000	Validation accuracy: 75.000	
Train Step: 14000	Train Loss: 0.117	Train Accuracy: 96.000
Validation step: 14000	Validation accuracy: 75.000	
Train Step: 15000	Train Loss: 0.071	Train Accuracy: 98.000
Validation step: 15000	Validation accuracy: 75.000	
Train Step: 16000	Train Loss: 0.034	Train Accuracy: 98.000
Validation step: 16000	Validation accuracy: 75.000	
Train Step: 17000	Train Loss: 0.055	Train Accuracy: 98.000
Validation step: 17000	Validation accuracy: 74.000	
Train Step: 18000	Train Loss: 0.050	Train Accuracy: 96.000
Validation step: 18000	Validation accuracy: 74.000	
Train Step: 19000	Train Loss: 0.128	Train Accuracy: 94.000
Validation step: 19000	Validation accuracy: 74.000	

```

In [0]: model1.eval()
        correct = 0
        for images, labels in testloader:
            with torch.no_grad(): # so that computation graph history is not stored
                images, labels = images.cuda(), labels.cuda() # send tensors to GPU
                outputs = model1(images)
                predictions = outputs.data.max(1)[1]
                correct += predictions.eq(labels.data).sum()

        print('Test set accuracy: {:.2f}%'.format(100.0 * correct / len(testloader.dataset)))

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:21: UserWarning: Implicit dimensi

Test set accuracy: 75.00%

```
In [0]: import torchvision
import torchvision.transforms as transforms
testloader = torch.utils.data.DataLoader(testset, batch_size=10,
                                         shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

import matplotlib.pyplot as plt
import numpy as np

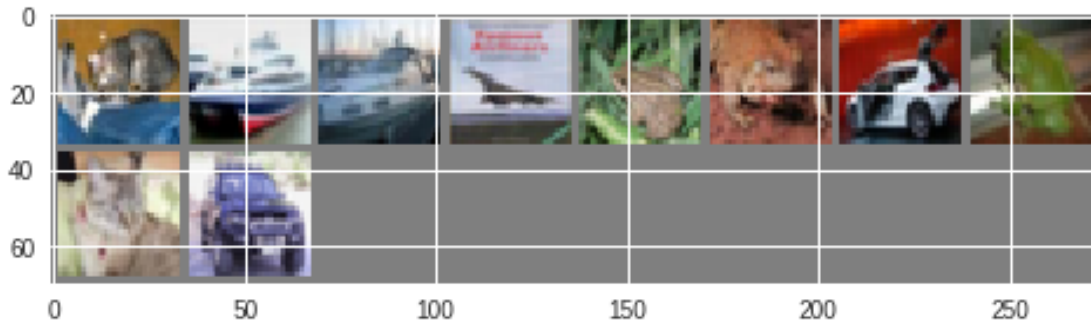
# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(testloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
images, labels = images.cuda(), labels.cuda() # send tensors to GPU
outputs = model1(images)
predictions_T = outputs.data.max(1)[1]
correct += predictions_T.eq(labels.data).sum()

# print labels
print("testlabels: " + ' '.join('%5s' % classes[labels.data[j]] for j in range(10)))
print("predlabels: " + ' '.join('%5s' % classes[predictions_T[j]] for j in range(10)))
```



```
testlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
predlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:21: UserWarning: Implicit dimension

```
In [0]: import torch.nn as nn
import torch.nn.functional as F

#model2 changes the size of kernel from 5x5 to 3x3
#model2 uses nn.Sequential mode to include a ReLu activation and a Dropout in the first
#model 2 improve valid_acc and test_acc to 76%
class CIFAR10_2(nn.Module):
    def __init__(self):
        super(CIFAR10_2, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.Dropout(p = 0.2),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.layer2 = nn.Sequential(
            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.Dropout(p = 0.2),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.drop_out = nn.Dropout()
        self.fc1 = nn.Linear(64*8*8, 1024)
        self.fc2 = nn.Linear(1024, 10)
    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(-1, 64*8*8) # reshape before sending to fc layer
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training) # default p=0.5
```

```

        x = self.fc2(x)
        return F.log_softmax(x)

model2 = CIFAR10_2()

In [0]: print(model2)

CIFAR10_2(
  (layer1): Sequential(
    (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Dropout(p=0.2)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (layer2): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Dropout(p=0.2)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (drop_out): Dropout(p=0.5)
  (fc1): Linear(in_features=4096, out_features=1024, bias=True)
  (fc2): Linear(in_features=1024, out_features=10, bias=True)
)

In [0]: import torch.nn.functional as F
        loss = F.nll_loss(outputs, labels)

```

1 New Section

```

In [0]: import torch.optim as optim

        criterion = nn.CrossEntropyLoss()
        optimizer = optim.SGD(model2.parameters(), lr=0.01, momentum=0.9)

In [0]: batch_size = 50
        # send model to GPU

        #train_loss, train_accu, valid_accu = [], [], []

        for epoch in range(40):
            train(model2, train_loader, criterion, optimizer, epoch)
            validation(model2, valid_loader, epoch)

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:29: UserWarning: Implicit dimensi

```

Train Step: 0	Train Loss: 2.305	Train Accuracy: 6.000
Validation step: 0	Validation accuracy: 52.000	
Train Step: 1000	Train Loss: 1.349	Train Accuracy: 52.000
Validation step: 1000	Validation accuracy: 65.000	
Train Step: 2000	Train Loss: 0.881	Train Accuracy: 72.000
Validation step: 2000	Validation accuracy: 68.000	
Train Step: 3000	Train Loss: 0.964	Train Accuracy: 68.000
Validation step: 3000	Validation accuracy: 72.000	
Train Step: 4000	Train Loss: 0.547	Train Accuracy: 86.000
Validation step: 4000	Validation accuracy: 72.000	
Train Step: 5000	Train Loss: 0.594	Train Accuracy: 80.000
Validation step: 5000	Validation accuracy: 72.000	
Train Step: 6000	Train Loss: 0.621	Train Accuracy: 74.000
Validation step: 6000	Validation accuracy: 73.000	
Train Step: 7000	Train Loss: 0.338	Train Accuracy: 86.000
Validation step: 7000	Validation accuracy: 74.000	
Train Step: 8000	Train Loss: 0.385	Train Accuracy: 88.000
Validation step: 8000	Validation accuracy: 74.000	
Train Step: 9000	Train Loss: 0.416	Train Accuracy: 84.000
Validation step: 9000	Validation accuracy: 75.000	
Train Step: 10000	Train Loss: 0.294	Train Accuracy: 92.000
Validation step: 10000	Validation accuracy: 75.000	
Train Step: 11000	Train Loss: 0.150	Train Accuracy: 94.000
Validation step: 11000	Validation accuracy: 76.000	
Train Step: 12000	Train Loss: 0.167	Train Accuracy: 96.000
Validation step: 12000	Validation accuracy: 75.000	
Train Step: 13000	Train Loss: 0.172	Train Accuracy: 94.000
Validation step: 13000	Validation accuracy: 75.000	
Train Step: 14000	Train Loss: 0.140	Train Accuracy: 94.000
Validation step: 14000	Validation accuracy: 75.000	
Train Step: 15000	Train Loss: 0.056	Train Accuracy: 98.000
Validation step: 15000	Validation accuracy: 75.000	
Train Step: 16000	Train Loss: 0.293	Train Accuracy: 90.000
Validation step: 16000	Validation accuracy: 75.000	
Train Step: 17000	Train Loss: 0.158	Train Accuracy: 94.000
Validation step: 17000	Validation accuracy: 76.000	
Train Step: 18000	Train Loss: 0.095	Train Accuracy: 96.000
Validation step: 18000	Validation accuracy: 75.000	
Train Step: 19000	Train Loss: 0.149	Train Accuracy: 94.000
Validation step: 19000	Validation accuracy: 76.000	
Train Step: 20000	Train Loss: 0.147	Train Accuracy: 98.000
Validation step: 20000	Validation accuracy: 76.000	
Train Step: 21000	Train Loss: 0.123	Train Accuracy: 96.000
Validation step: 21000	Validation accuracy: 75.000	
Train Step: 22000	Train Loss: 0.350	Train Accuracy: 94.000
Validation step: 22000	Validation accuracy: 76.000	
Train Step: 23000	Train Loss: 0.219	Train Accuracy: 96.000
Validation step: 23000	Validation accuracy: 76.000	

Train Step: 24000	Train Loss: 0.044	Train Accuracy: 100.000
Validation step: 24000	Validation accuracy: 76.000	
Train Step: 25000	Train Loss: 0.188	Train Accuracy: 94.000
Validation step: 25000	Validation accuracy: 75.000	
Train Step: 26000	Train Loss: 0.074	Train Accuracy: 98.000
Validation step: 26000	Validation accuracy: 76.000	
Train Step: 27000	Train Loss: 0.026	Train Accuracy: 98.000
Validation step: 27000	Validation accuracy: 75.000	
Train Step: 28000	Train Loss: 0.106	Train Accuracy: 96.000
Validation step: 28000	Validation accuracy: 75.000	
Train Step: 29000	Train Loss: 0.230	Train Accuracy: 94.000
Validation step: 29000	Validation accuracy: 76.000	
Train Step: 30000	Train Loss: 0.155	Train Accuracy: 94.000
Validation step: 30000	Validation accuracy: 76.000	
Train Step: 31000	Train Loss: 0.029	Train Accuracy: 100.000
Validation step: 31000	Validation accuracy: 76.000	
Train Step: 32000	Train Loss: 0.033	Train Accuracy: 100.000
Validation step: 32000	Validation accuracy: 76.000	
Train Step: 33000	Train Loss: 0.039	Train Accuracy: 100.000
Validation step: 33000	Validation accuracy: 76.000	
Train Step: 34000	Train Loss: 0.020	Train Accuracy: 100.000
Validation step: 34000	Validation accuracy: 75.000	
Train Step: 35000	Train Loss: 0.041	Train Accuracy: 96.000
Validation step: 35000	Validation accuracy: 75.000	
Train Step: 36000	Train Loss: 0.024	Train Accuracy: 100.000
Validation step: 36000	Validation accuracy: 76.000	
Train Step: 37000	Train Loss: 0.104	Train Accuracy: 96.000
Validation step: 37000	Validation accuracy: 75.000	
Train Step: 38000	Train Loss: 0.095	Train Accuracy: 94.000
Validation step: 38000	Validation accuracy: 76.000	
Train Step: 39000	Train Loss: 0.274	Train Accuracy: 90.000
Validation step: 39000	Validation accuracy: 76.000	

```
In [0]: model2.eval()
        correct = 0
        for images, labels in testloader:
            with torch.no_grad(): # so that computation graph history is not stored
                images, labels = images.cuda(), labels.cuda() # send tensors to GPU
                outputs = model2(images)
                predictions = outputs.data.max(1)[1]
                correct += predictions.eq(labels.data).sum()

        print('Test set accuracy: {:.2f}%'.format(100.0 * correct / len(testloader.dataset)))
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:29: UserWarning: Implicit dimensions

Test set accuracy: 76.00%

```

In [0]: import torchvision
import torchvision.transforms as transforms
testloader = torch.utils.data.DataLoader(testset, batch_size=10,
                                         shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image


def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(testloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
images, labels = images.cuda(), labels.cuda() # send tensors to GPU
outputs = model2(images)
predictions_T = outputs.data.max(1)[1]
correct += predictions_T.eq(labels.data).sum()

# print labels
print("testlabels: " + ' '.join('%5s' % classes[labels.data[j]] for j in range(10)))
print("predlabels: " + ' '.join('%5s' % classes[predictions_T[j]] for j in range(10)))

```



```
testlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
predlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:29: UserWarning: Implicit dimension

```
In [0]: import torch.nn as nn
import torch.nn.functional as F
# Model3 adds batchNorm2d layer after ReLu activation in the first two convolutional l
# Model3 add fc3 after fc2 in fully-connected layers
# Model3 still incooperates dropout layers in both convolutional and fully-connected l
# Model3 has improved valid_acc and test_acc to 78%

class CIFAR10_3(nn.Module):
    def __init__(self):
        super(CIFAR10_3, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.BatchNorm2d(32),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Dropout(p=0.2))
        self.layer2 = nn.Sequential(
            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.BatchNorm2d(64),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Dropout(p=0.2))
        self.drop_out = nn.Dropout()
        self.fc1 = nn.Linear(64*8*8, 1024)
        self.fc2 = nn.Linear(1024, 100)
        self.fc3 = nn.Linear(100, 10)
    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(-1, 64*8*8) # reshape before sending to fc layer
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = F.relu(self.fc2(x))
        x = F.dropout(x, training=self.training) # default p=0.5
        x = self.fc3(x)
        return F.log_softmax(x)

model3 = CIFAR10_3()
```

```

In [0]: print(model3)

CIFAR10_3(
  (layer1): Sequential(
    (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (4): Dropout(p=0.2)
  )
  (layer2): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (4): Dropout(p=0.2)
  )
  (drop_out): Dropout(p=0.5)
  (fc1): Linear(in_features=4096, out_features=1024, bias=True)
  (fc2): Linear(in_features=1024, out_features=100, bias=True)
  (fc3): Linear(in_features=100, out_features=10, bias=True)
)

In [0]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model3.parameters(), lr=0.01, momentum=0.9)

In [0]: batch_size = 50
        # send model to GPU

        #train_loss, train_accu, valid_accu = [], [], []

        for epoch in range(40):
            train(model3, train_loader, criterion, optimizer, epoch)
            validation(model3, valid_loader, epoch)

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:36: UserWarning: Implicit dimensi

```

Train Step: 0	Train Loss: 2.317	Train Accuracy: 8.000
Validation step: 0	Validation accuracy: 56.000	
Train Step: 1000	Train Loss: 1.307	Train Accuracy: 48.000
Validation step: 1000	Validation accuracy: 65.000	
Train Step: 2000	Train Loss: 1.079	Train Accuracy: 60.000
Validation step: 2000	Validation accuracy: 70.000	
Train Step: 3000	Train Loss: 0.750	Train Accuracy: 78.000
Validation step: 3000	Validation accuracy: 71.000	

Train Step: 4000	Train Loss: 1.115	Train Accuracy: 66.000
Validation step: 4000	Validation accuracy: 72.000	
Train Step: 5000	Train Loss: 0.736	Train Accuracy: 72.000
Validation step: 5000	Validation accuracy: 74.000	
Train Step: 6000	Train Loss: 0.775	Train Accuracy: 70.000
Validation step: 6000	Validation accuracy: 74.000	
Train Step: 7000	Train Loss: 0.611	Train Accuracy: 74.000
Validation step: 7000	Validation accuracy: 75.000	
Train Step: 8000	Train Loss: 0.651	Train Accuracy: 72.000
Validation step: 8000	Validation accuracy: 75.000	
Train Step: 9000	Train Loss: 0.675	Train Accuracy: 76.000
Validation step: 9000	Validation accuracy: 76.000	
Train Step: 10000	Train Loss: 0.534	Train Accuracy: 84.000
Validation step: 10000	Validation accuracy: 76.000	
Train Step: 11000	Train Loss: 0.636	Train Accuracy: 68.000
Validation step: 11000	Validation accuracy: 76.000	
Train Step: 12000	Train Loss: 0.480	Train Accuracy: 86.000
Validation step: 12000	Validation accuracy: 77.000	
Train Step: 13000	Train Loss: 0.484	Train Accuracy: 82.000
Validation step: 13000	Validation accuracy: 75.000	
Train Step: 14000	Train Loss: 0.376	Train Accuracy: 90.000
Validation step: 14000	Validation accuracy: 76.000	
Train Step: 15000	Train Loss: 0.351	Train Accuracy: 92.000
Validation step: 15000	Validation accuracy: 77.000	
Train Step: 16000	Train Loss: 0.318	Train Accuracy: 86.000
Validation step: 16000	Validation accuracy: 77.000	
Train Step: 17000	Train Loss: 0.398	Train Accuracy: 84.000
Validation step: 17000	Validation accuracy: 77.000	
Train Step: 18000	Train Loss: 0.402	Train Accuracy: 88.000
Validation step: 18000	Validation accuracy: 78.000	
Train Step: 19000	Train Loss: 0.201	Train Accuracy: 94.000
Validation step: 19000	Validation accuracy: 77.000	
Train Step: 20000	Train Loss: 0.243	Train Accuracy: 96.000
Validation step: 20000	Validation accuracy: 77.000	
Train Step: 21000	Train Loss: 0.229	Train Accuracy: 92.000
Validation step: 21000	Validation accuracy: 78.000	
Train Step: 22000	Train Loss: 0.170	Train Accuracy: 94.000
Validation step: 22000	Validation accuracy: 78.000	
Train Step: 23000	Train Loss: 0.382	Train Accuracy: 84.000
Validation step: 23000	Validation accuracy: 77.000	
Train Step: 24000	Train Loss: 0.401	Train Accuracy: 88.000
Validation step: 24000	Validation accuracy: 78.000	
Train Step: 25000	Train Loss: 0.279	Train Accuracy: 86.000
Validation step: 25000	Validation accuracy: 78.000	
Train Step: 26000	Train Loss: 0.161	Train Accuracy: 94.000
Validation step: 26000	Validation accuracy: 78.000	
Train Step: 27000	Train Loss: 0.296	Train Accuracy: 92.000
Validation step: 27000	Validation accuracy: 77.000	

Train Step: 28000	Train Loss: 0.244	Train Accuracy: 94.000
Validation step: 28000	Validation accuracy: 78.000	
Train Step: 29000	Train Loss: 0.171	Train Accuracy: 92.000
Validation step: 29000	Validation accuracy: 77.000	
Train Step: 30000	Train Loss: 0.230	Train Accuracy: 94.000
Validation step: 30000	Validation accuracy: 78.000	
Train Step: 31000	Train Loss: 0.254	Train Accuracy: 88.000
Validation step: 31000	Validation accuracy: 78.000	
Train Step: 32000	Train Loss: 0.210	Train Accuracy: 96.000
Validation step: 32000	Validation accuracy: 78.000	
Train Step: 33000	Train Loss: 0.304	Train Accuracy: 92.000
Validation step: 33000	Validation accuracy: 77.000	
Train Step: 34000	Train Loss: 0.120	Train Accuracy: 96.000
Validation step: 34000	Validation accuracy: 77.000	
Train Step: 35000	Train Loss: 0.225	Train Accuracy: 94.000
Validation step: 35000	Validation accuracy: 77.000	
Train Step: 36000	Train Loss: 0.082	Train Accuracy: 96.000
Validation step: 36000	Validation accuracy: 78.000	
Train Step: 37000	Train Loss: 0.240	Train Accuracy: 90.000
Validation step: 37000	Validation accuracy: 78.000	
Train Step: 38000	Train Loss: 0.101	Train Accuracy: 96.000
Validation step: 38000	Validation accuracy: 78.000	
Train Step: 39000	Train Loss: 0.239	Train Accuracy: 90.000
Validation step: 39000	Validation accuracy: 78.000	

```
In [0]: model3.eval()
        correct = 0
        for images, labels in testloader:
            with torch.no_grad(): # so that computation graph history is not stored
                images, labels = images.cuda(), labels.cuda() # send tensors to GPU
                outputs = model3(images)
                predictions = outputs.data.max(1)[1]
                correct += predictions.eq(labels.data).sum()

        print('Test set accuracy: {:.2f}%'.format(100.0 * correct / len(testloader.dataset)))
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:36: UserWarning: Implicit dimension

Test set accuracy: 78.00%

```
In [0]: import torchvision
        import torchvision.transforms as transforms
        testloader = torch.utils.data.DataLoader(testset, batch_size=10,
                                                    shuffle=False, num_workers=2)

        classes = ('plane', 'car', 'bird', 'cat',
```

```

        'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

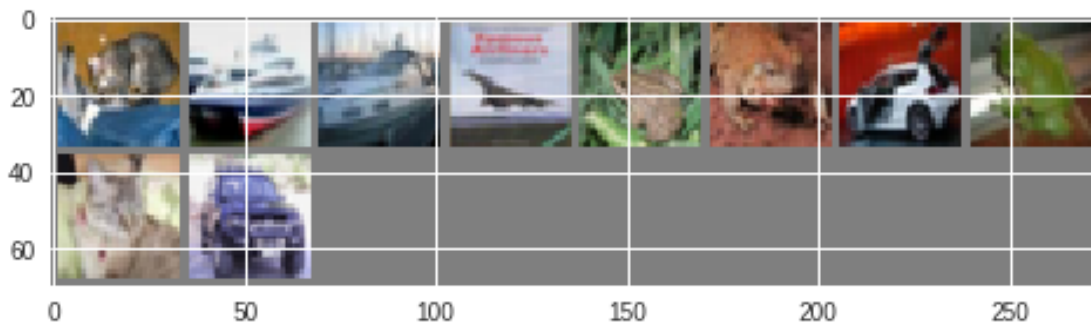
def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(testloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
images, labels = images.cuda(), labels.cuda() # send tensors to GPU
outputs = model3(images)
predictions_T = outputs.data.max(1)[1]
correct += predictions_T.eq(labels.data).sum()

# print labels
print("testlabels: " + ' '.join('%5s' % classes[labels.data[j]] for j in range(10)))
print("predlabels: " + ' '.join('%5s' % classes[predictions_T[j]] for j in range(10)))

```



```

testlabels:  cat  ship  ship plane  frog  frog  car  frog  cat  car
predlabels:  cat  ship  ship plane  cat  frog  car  frog  cat  car

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:36: UserWarning: Implicit dimension

```

In [0]: import torch.nn as nn
import torch.nn.functional as F
# Model3 adds batchNorm2d layer after ReLu activation in the first two convolutional l
# Model3 add fc3 after fc2 in fully-connected layers
# Model3 still incooperates dropout layers in both convolutional and fully-connected l
# Model3 has improved valid_acc and test_acc to 79%
# Model4 keeps kernel size at 5x5

class CIFAR10_4(nn.Module):
    def __init__(self):
        super(CIFAR10_4, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=5, stride=1, padding=2),
            nn.ReLU(),
            nn.BatchNorm2d(32),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Dropout(p=0.2))
        self.layer2 = nn.Sequential(
            nn.Conv2d(32, 64, kernel_size=5, stride=1, padding=2),
            nn.ReLU(),
            nn.BatchNorm2d(64),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Dropout(p=0.2))
        self.drop_out = nn.Dropout()
        self.fc1 = nn.Linear(64*8*8, 1024)
        self.fc2 = nn.Linear(1024, 100)
        self.fc3 = nn.Linear(100, 10)
    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(-1, 64*8*8) # reshape before sending to fc layer
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = F.relu(self.fc2(x))
        x = F.dropout(x, training=self.training) # default p=0.5
        x = self.fc3(x)
        return F.log_softmax(x)
model4 = CIFAR10_4()

```

```

In [0]: print(model4)

```

```

CIFAR10_4(
  (layer1): Sequential(
    (0): Conv2d(3, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

```



```

        (4): Dropout(p=0.2)
    )
    (layer2): Sequential(
      (0): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (1): ReLU()
      (2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
      (4): Dropout(p=0.2)
    )
    (drop_out): Dropout(p=0.5)
    (fc1): Linear(in_features=4096, out_features=1024, bias=True)
    (fc2): Linear(in_features=1024, out_features=100, bias=True)
    (fc3): Linear(in_features=100, out_features=10, bias=True)
  )

```

```

In [0]: import torch.nn.functional as F
        loss = F.nll_loss(outputs, labels)

```

```

In [0]: import torch.optim as optim

        criterion = nn.CrossEntropyLoss()
        optimizer = optim.SGD(model4.parameters(), lr=0.01, momentum=0.9)

```

```

In [0]: batch_size = 50
        # send model to GPU

        #train_loss, train_accu, valid_accu = [], [], []

        for epoch in range(40):
            train(model4, train_loader, criterion, optimizer, epoch)
            validation(model4, valid_loader, epoch)

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:34: UserWarning: Implicit dimension

Train Step: 0	Train Loss: 2.322	Train Accuracy: 10.000
Validation step: 0	Validation accuracy: 57.000	
Train Step: 1000	Train Loss: 1.483	Train Accuracy: 56.000
Validation step: 1000	Validation accuracy: 65.000	
Train Step: 2000	Train Loss: 1.118	Train Accuracy: 62.000
Validation step: 2000	Validation accuracy: 68.000	
Train Step: 3000	Train Loss: 0.888	Train Accuracy: 68.000
Validation step: 3000	Validation accuracy: 73.000	
Train Step: 4000	Train Loss: 0.867	Train Accuracy: 74.000
Validation step: 4000	Validation accuracy: 73.000	
Train Step: 5000	Train Loss: 0.830	Train Accuracy: 76.000
Validation step: 5000	Validation accuracy: 75.000	
Train Step: 6000	Train Loss: 0.744	Train Accuracy: 66.000

Validation step: 6000	Validation accuracy: 75.000
Train Step: 7000	Train Loss: 0.697 Train Accuracy: 78.000
Validation step: 7000	Validation accuracy: 75.000
Train Step: 8000	Train Loss: 0.769 Train Accuracy: 70.000
Validation step: 8000	Validation accuracy: 75.000
Train Step: 9000	Train Loss: 0.771 Train Accuracy: 80.000
Validation step: 9000	Validation accuracy: 76.000
Train Step: 10000	Train Loss: 0.453 Train Accuracy: 88.000
Validation step: 10000	Validation accuracy: 76.000
Train Step: 11000	Train Loss: 0.646 Train Accuracy: 76.000
Validation step: 11000	Validation accuracy: 77.000
Train Step: 12000	Train Loss: 0.392 Train Accuracy: 84.000
Validation step: 12000	Validation accuracy: 77.000
Train Step: 13000	Train Loss: 0.302 Train Accuracy: 90.000
Validation step: 13000	Validation accuracy: 77.000
Train Step: 14000	Train Loss: 0.458 Train Accuracy: 80.000
Validation step: 14000	Validation accuracy: 76.000
Train Step: 15000	Train Loss: 0.221 Train Accuracy: 90.000
Validation step: 15000	Validation accuracy: 77.000
Train Step: 16000	Train Loss: 0.243 Train Accuracy: 92.000
Validation step: 16000	Validation accuracy: 77.000
Train Step: 17000	Train Loss: 0.453 Train Accuracy: 88.000
Validation step: 17000	Validation accuracy: 77.000
Train Step: 18000	Train Loss: 0.378 Train Accuracy: 84.000
Validation step: 18000	Validation accuracy: 77.000
Train Step: 19000	Train Loss: 0.237 Train Accuracy: 90.000
Validation step: 19000	Validation accuracy: 78.000
Train Step: 20000	Train Loss: 0.326 Train Accuracy: 92.000
Validation step: 20000	Validation accuracy: 78.000
Train Step: 21000	Train Loss: 0.281 Train Accuracy: 94.000
Validation step: 21000	Validation accuracy: 78.000
Train Step: 22000	Train Loss: 0.092 Train Accuracy: 98.000
Validation step: 22000	Validation accuracy: 78.000
Train Step: 23000	Train Loss: 0.094 Train Accuracy: 96.000
Validation step: 23000	Validation accuracy: 78.000
Train Step: 24000	Train Loss: 0.229 Train Accuracy: 94.000
Validation step: 24000	Validation accuracy: 79.000
Train Step: 25000	Train Loss: 0.289 Train Accuracy: 88.000
Validation step: 25000	Validation accuracy: 78.000
Train Step: 26000	Train Loss: 0.200 Train Accuracy: 92.000
Validation step: 26000	Validation accuracy: 79.000
Train Step: 27000	Train Loss: 0.117 Train Accuracy: 96.000
Validation step: 27000	Validation accuracy: 79.000
Train Step: 28000	Train Loss: 0.170 Train Accuracy: 94.000
Validation step: 28000	Validation accuracy: 78.000
Train Step: 29000	Train Loss: 0.028 Train Accuracy: 100.000
Validation step: 29000	Validation accuracy: 78.000
Train Step: 30000	Train Loss: 0.230 Train Accuracy: 92.000

Validation step: 30000	Validation accuracy: 78.000
Train Step: 31000	Train Loss: 0.154 Train Accuracy: 92.000
Validation step: 31000	Validation accuracy: 78.000
Train Step: 32000	Train Loss: 0.134 Train Accuracy: 96.000
Validation step: 32000	Validation accuracy: 79.000
Train Step: 33000	Train Loss: 0.110 Train Accuracy: 94.000
Validation step: 33000	Validation accuracy: 79.000
Train Step: 34000	Train Loss: 0.097 Train Accuracy: 98.000
Validation step: 34000	Validation accuracy: 78.000
Train Step: 35000	Train Loss: 0.215 Train Accuracy: 94.000
Validation step: 35000	Validation accuracy: 78.000
Train Step: 36000	Train Loss: 0.042 Train Accuracy: 100.000
Validation step: 36000	Validation accuracy: 79.000
Train Step: 37000	Train Loss: 0.068 Train Accuracy: 96.000
Validation step: 37000	Validation accuracy: 78.000
Train Step: 38000	Train Loss: 0.090 Train Accuracy: 96.000
Validation step: 38000	Validation accuracy: 79.000
Train Step: 39000	Train Loss: 0.058 Train Accuracy: 98.000
Validation step: 39000	Validation accuracy: 79.000

```
In [0]: model4.eval()
        correct = 0
        for images, labels in testloader:
            with torch.no_grad(): # so that computation graph history is not stored
                images, labels = images.cuda(), labels.cuda() # send tensors to GPU
                outputs = model4(images)
                predictions = outputs.data.max(1)[1]
                correct += predictions.eq(labels.data).sum()

        print('Test set accuracy: {:.2f}%'.format(100.0 * correct / len(testloader.dataset)))
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:34: UserWarning: Implicit dimension

Test set accuracy: 79.00%

```
In [ ]: import torchvision
        import torchvision.transforms as transforms
        testloader = torch.utils.data.DataLoader(testset, batch_size=10,
                                                    shuffle=False, num_workers=2)

        classes = ('plane', 'car', 'bird', 'cat',
                    'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

        import matplotlib.pyplot as plt
        import numpy as np
```

```

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(testloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
images, labels = images.cuda(), labels.cuda() # send tensors to GPU
outputs = model4(images)
predictions_T = outputs.data.max(1)[1]
correct += predictions_T.eq(labels.data).sum()

# print labels
print("testlabels: " + ' '.join('%5s' % classes[labels.data[j]] for j in range(10)))
print("predlabels: " + ' '.join('%5s' % classes[predictions_T[j]] for j in range(10)))

```

Discussion for Part 2:

I start with my base line model of two convolutional layers: each of kernel size 5, the first layer has input channel 3 and output channel 32 and second layer has input channel 32 and output channel 64. Following each of these convolutional layers, I have maxpool layer of 2x2. After passing data through these two convolution layers, I then pass feature maps(tensors) to two fcs: fc1 (64x8x8, 1024), fc2(1024, 10). For fc1, there is relu activation while fc2 does not. After linear portion of fc2, data are passed through softmax layer. The cost function I chose was crossentropyloss and optimization I used was gradient descent . With this base line model cNN, I got 73% for validation and test accuracy. One problem I noticed was that the model was overfitted by training data during training. I had to stop training after 8000 iterations as training accuracy reach 100%.

The first thing I did to modify this model was to add a dropout layer after fc2. This was my model1. This obviously have controlled overfitting problem better. I was able to run more iterations. In the end, I was able to get 74% for validation accuracy and 75% for test accuracy after adding this dropout layer.

I also tried a number of other modifications: including change the kernel size, add more fcs, and add batchnorm layer and relu layer in convolutional layers of my model. Here are what I observed:

modifications	valid_acc	test_acc
model_b None	73%	73%
model1 dropout after fc1 (this dropout is included in model2,3,4)	75%	74%

model2 kernel 3x3, add relu and dropout 76% 76%
in convolutional layers
model3 kernel 3x3, add batch, relu and dropout 78% 78%
in convolutional layers, add fc3 after fc2 and dropout
model4 kernel 5x5(the same with model_b), 79% 79%
add batch, relu and dropout
in convolutional layers, add fc3 after fc2 and dropout

My conclusion: Dropout and batchnorm technique can improve generalization of CNN model and avoid overfitting of training data. In addition, altering architecture of CNN can also affect performance of CNN model. A relatively larger size of kernel seems to work better than smaller size kernel and it captures more information in feature maps.